

# Programming for high performance

- **Optimizing the performance of parallel programs is an iterative process of refining choices for decomposition, assignment, and orchestration...**
- **Key goals (that are at odds with each other)**
  - **Balance workload onto available execution resources**
  - **Reduce communication (to avoid stalls)**
  - **Reduce extra work (overhead) performed to increase parallelism, manage assignment, reduce communication, etc.**
- **We are going to talk about a rich space of techniques**

**TIP #1: Always implement the simplest solution first, then measure performance to determine if you need to do better.**

**“My solution scales” = your code scales as much as you need it to.**

**(if you anticipate only running low-core count machines, it may be unnecessary to implement a complex approach that creates and hundreds or thousands of pieces of independent work)**

# Balancing the workload

**Ideally: all processors are computing all the time during program execution  
(they are computing simultaneously, and they finish their portion of the work at the same time)**




**Recall Amdahl's Law:  
Only small amount of load imbalance can  
significantly bound maximum speedup**

**P4 does 20% more work → P4 takes 20% longer to complete  
→ 20% of parallel program's  
runtime is serial execution**

**(work in serialized section here is about 5% of the work of the whole program:  
 $S=.05$  in Amdahl's law equation)**

# Static assignment


- Assignment of work to threads is pre-determined
  - Not necessarily determined at compile-time (assignment algorithm may depend on runtime parameters such as input data size, number of threads, etc.)
- Recall solver example: assign equal number of grid cells (work) to each thread (worker)
  - We discussed two static assignments of work to workers (blocked and interleaved)



- Good properties of static assignment: simple, essentially zero runtime overhead (in this example: extra work to implement assignment is a little bit of indexing math)

# When is static assignment applicable?

- When the cost (execution time) of work and the amount of work is predictable (so the programmer can work out a good assignment in advance)
- Simplest example: it is known up front that all work has the same cost




**In the example above:**

**There are 12 tasks, and it is known that each have the same cost.**

**Assignment solution: statically assign three tasks to each of the four processors.**

# When is static assignment applicable?


- When work is predictable, but not all jobs have same cost (see example below)
- When statistics about execution time are known (e.g., same cost on average)



Jobs have unequal, but known cost: assign to processors to ensure overall good load balance

# “Semi-static” assignment

- Cost of work is predictable for near-term future
  - Idea: recent past good predictor of near future
- Application periodically profiles application and re-adjusts assignment
  - Assignment is “static” for the interval between re-adjustments



**Particle simulation:**

Redistribute particles as they move over course of simulation  
(if motion is slow, redistribution need not occur often)




Image credit: <http://typhon.sourceforge.net/spip/spip.php?article22>

**Adaptive mesh:**

Mesh is changed as object moves or flow over object changes, but changes occur slowly (color indicates assignment of parts of mesh to processors)



# Dynamic assignment

Program determines assignment dynamically at runtime to ensure a well distributed load. (The execution time of tasks, or the total number of tasks, is unpredictable.)

**Sequential program**  
(independent loop iterations)

```
int N = 1024;
int* x = new int[N];
bool* prime = new bool[N];

// initialize elements of x here

for (int i=0; i<N; i++)
{
    // unknown execution time
    is_prime[i] = test_primalty(x[i]);
}
```

**Parallel program**  
(SPMD execution by multiple threads,  
shared address space model)

```
int N = 1024;
// assume allocations are only executed by 1 thread
int* x = new int[N];
bool* is_prime = new bool[N];

// initialize elements of x here


LOCK counter_lock;
int counter = 0;    // shared variable

while (1) {
    int i;
    lock(counter_lock);
    i = counter++;
    unlock(counter_lock);
    atomic_incr(counter);
    if (i >= N)
        break;
    is_prime[i] = test_primalty(x[i]);
}
```




# Dynamic assignment using a work queue

**Sub-problems**  
(a.k.a. “tasks”, “work”)



**Shared work queue: a list of work to do**  
(for now, let's assume each piece of work is independent)



**Worker threads:**  
Pull data from shared work queue  
Push new work to queue as it is created

# What constitutes a piece of work?

## What is a potential problem with this implementation?

```
const int N = 1024;
// assume allocations are only executed by 1 thread
float* x = new float[N];
bool* prime = new bool[N];

// initialize elements of x here

LOCK counter_lock;
int counter = 0;


while (1) {
    int i;
    lock(counter_lock);
    i = counter++;
    unlock(counter_lock);
    if (i >= N)
        break;
    is_prime[i] = test_primalty(x[i]);
}
```

Time in task 0 ————— I

Time in critical section ————— I

This is overhead that  
does not exist in serial  
program

And.. it's serial execution  
(recall Amdahl's Law)



Fine granularity partitioning: 1 “task” = 1 element

Likely good workload balance (many small tasks)

Potential for high synchronization cost  
(serialization at critical section)

## So... IS IT a problem?

# Increasing task granularity


```
const int N = 1024;
const int GRANULARITY = 10;
// assume allocations are only executed by 1 thread

float* x = new float[N];
bool* prime = new bool[N];

// initialize elements of x here

LOCK counter_lock;
int counter = 0;

while (1) {
    int i;
    lock(counter_lock);
    i = counter;
    counter += GRANULARITY;
    unlock(counter_lock);
    if (i >= N)
        break;
    int end = min(i + GRANULARITY, N);
    for (int j=i; j<end; j++)
        is_prime[j] = test_primalty(x[j]);
}
```



**Coarse granularity partitioning: 1 “task” = 10 elements**  
**Decreased synchronization cost**  
**(Critical section entered 10 times less)**

# Choosing task size


- **Useful to have many more tasks\* than processors**  
(many small tasks enables good workload balance via dynamic assignment)
  - Motivates small granularity tasks
- **But want as few tasks as possible to minimize overhead of managing the assignment**
  - Motivates large granularity tasks
- **Ideal granularity depends on many factors**  
(Common theme in this course: must know your workload, and your machine)

\* I had to pick a term for a piece of work, a sub-problem, etc.

# Smarter task scheduling


Consider dynamic scheduling via a shared work queue

What happens if the system assigns these tasks to workers in left-to-right order?



# Smarter task scheduling

**What happens if scheduler runs the long task last? Potential for load imbalance!**




**One possible solution to imbalance problem:**

**Divide work into a larger number of smaller tasks**

- Hopefully “long pole” gets shorter relative to overall execution time
- May increase synchronization overhead
- May not be possible (perhaps long task is fundamentally sequential)

# Smarter task scheduling

Schedule long task first to reduce “slop” at end of computation



Another solution: smarter scheduling

Schedule long tasks first

- Thread performing long task performs fewer overall tasks, but approximately the same amount of work as the other threads.
- Requires some knowledge of workload (some predictability of cost)